

# COMS 4721: Machine Learning for Data Science

## Lecture 10, 2/16/2021

Prof. John Paisley

Department of Electrical Engineering

Columbia University

# FEATURE EXPANSIONS

# FEATURE EXPANSIONS

**Feature expansions** (also called **basis expansions**) are names given to a technique we've already discussed and made use of.

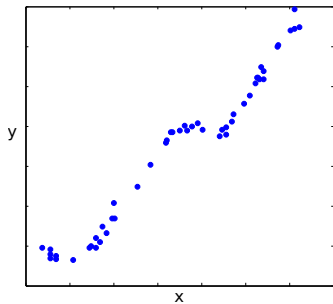
**Problem:** A linear model on the original feature space  $x \in \mathbb{R}^d$  doesn't work.

**Solution:** Map the features to another space  $\phi(x) \in \mathbb{R}^D$ , often where  $D > d$ , and do linear modeling there.

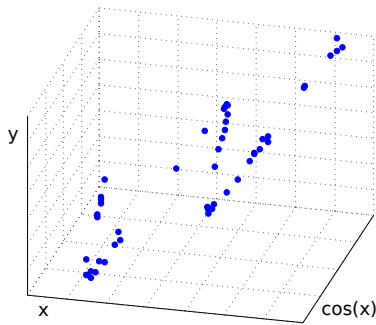
## Examples

- ▶ For polynomial regression on  $\mathbb{R}$ , we let  $\phi(x) = (x, x^2, \dots, x^p)^T$ .
- ▶ For one jump discontinuity, we could set  $\phi(x) = (x, \mathbb{1}\{x < a\})^T$ .

# MAPPING EXAMPLE FOR REGRESSION



(a) Data for linear regression



(b) Same data mapped to higher dimension

High-dimensional maps can transform the data so output is linear in inputs.


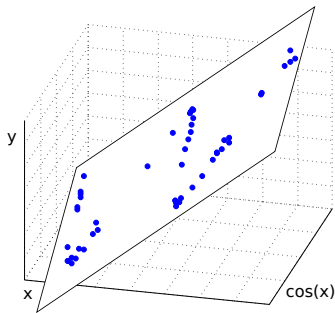
**Left:** Original  $x \in \mathbb{R}$  and response  $y$ .

**Right:**  $x$  mapped to  $\mathbb{R}^2$  using  $\phi(x) = (x, \cos x)^T$ .

# MAPPING EXAMPLE FOR REGRESSION

Using the mapping  $\phi(x) = (x, \cos x)^T$ , learn the linear regression model

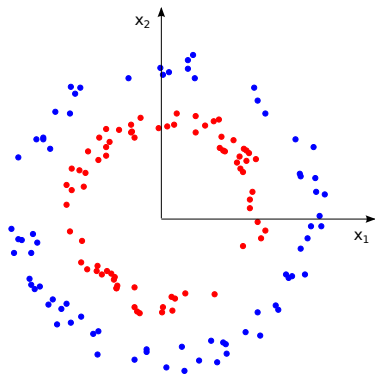
$$\begin{aligned}y &\approx w_0 + \phi(x)^T w \\ &\approx w_0 + w_1 x + w_2 \cos x.\end{aligned}$$




**Left:** Learn  $(w_0, w_1, w_2)$  to approximate data on the left with a plane.

**Right:** For each point  $x$ , map to  $\phi(x)$  and predict  $y$ . Plot as a function of  $x$ .

# MAPPING EXAMPLE FOR CLASSIFICATION



(e) Data for binary classification



(f) Same data mapped to higher dimension

High-dimensional maps can transform data so it becomes linearly separable.

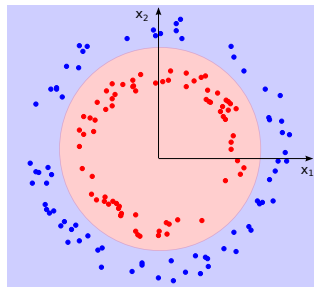

Left: Original data in  $\mathbb{R}^2$ .

Right: Data mapped to  $\mathbb{R}^3$  using  $\phi(x) = (x_1^2, x_1x_2, x_2^2)^T$ .

# MAPPING EXAMPLE FOR CLASSIFICATION

Using the mapping  $\phi(x) = (x_1^2, x_1x_2, x_2^2)^T$ , learn a linear classifier

$$\begin{aligned} y &= \text{sign}(w_0 + \phi(x)^T w) \\ &= \text{sign}(w_0 + w_1x_1^2 + w_2x_1x_2 + w_3x_2^2). \end{aligned}$$



**Left:** Learn  $(w_0, w_1, w_2, w_3)$  to linearly separate classes with hyperplane.

**Right:** For each point  $x$ , map to  $\phi(x)$  and classify. Color decision regions in  $\mathbb{R}^2$ .

# FEATURE EXPANSIONS AND DOT PRODUCTS

## What expansion should I use?

This is not obvious. The illustrations required knowledge about the data that we likely won't have (especially if it's in high dimensions).

One approach is to use the “kitchen sink”: If you can think of it, then use it. Select the useful features with an  $\ell_1$  penalty

$$w_{\ell_1} = \arg \min_w \sum_{i=1}^n f(y_i, \phi(x_i), w) + \lambda \|w\|_1.$$

We know that this will find a sparse subset of the dimensions of  $\phi(x)$  to use.

Often we only need to work with dot products  $\phi(x_i)^T \phi(x_j) \equiv K(x_i, x_j)$ . This is called a **kernel** and can produce some interesting results.



# KERNELS

# PERCEPTRON (SOME MOTIVATION)

## Perceptron classifier

Let  $x_i \in \mathbb{R}^{d+1}$  and  $y_i \in \{-1, +1\}$  for  $i = 1, \dots, n$  observations. We saw that the Perceptron constructs the hyperplane from data,

$$w = \sum_{i \in \mathcal{S}} y_i x_i, \quad (\text{assume } \eta = 1 \text{ and } \mathcal{S} \text{ has no duplicates})$$

where  $\mathcal{S}$  is the sequentially constructed set of misclassified examples.

## Predicting new data

We also discussed how we can predict the label  $y_0$  for a new observation  $x_0$ :

$$y_0 = \text{sign}(x_0^T w) = \text{sign} \left( \sum_{i \in \mathcal{S}} y_i x_0^T x_i \right)$$

We've taken feature expansions for granted, but we can explicitly write it as

$$y_0 = \text{sign}(\phi(x_0)^T w) = \text{sign} \left( \sum_{i \in \mathcal{S}} y_i \phi(x_0)^T \phi(x_i) \right)$$

We can represent the decision using dot products between data points.

# KERNELS

## Kernel definition

A kernel  $K(\cdot, \cdot) : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  is a symmetric function defined as follows:

**Definition:** If for any  $n$  points  $x_1, \dots, x_n \in \mathbb{R}^d$ , the  $n \times n$  matrix  $K$ , where  $K_{ij} = K(x_i, x_j)$ , is *positive semidefinite*, then  $K(\cdot, \cdot)$  is a “kernel.”

Intuitively, this means  $K$  satisfies the properties of a covariance matrix.

## Mercer's theorem

If the function  $K(\cdot, \cdot)$  satisfies the above properties, then there exists a mapping  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^D$  ( $D$  can equal  $\infty$ ) such that

$$K(x_i, x_j) = \phi(x_i)^T \phi(x_j).$$

If we first define  $\phi(\cdot)$  and then  $K$ , this is obvious. However, sometimes we first define  $K(\cdot, \cdot)$  and avoid ever using  $\phi(\cdot)$ .

# GAUSSIAN KERNEL (RADIAL BASIS FUNCTION)

The most popular kernel is the Gaussian kernel, also called the radial basis function (RBF),

$$K(x, x') = a \exp \left\{ -\frac{1}{b} \|x - x'\|^2 \right\}.$$

- ▶ This is a good, general-purpose kernel that usually works well.
- ▶ It takes into account proximity in  $\mathbb{R}^d$ . Things close together in space have larger value (as defined by kernel width  $b$ ).

In this case, the the mapping  $\phi(x)$  that produces the RBF kernel is *infinite dimensional* (it's a continuous function instead of a vector). Therefore

$$K(x, x') = \int \phi_t(x) \phi_t(x') dt.$$

- ▶  $\phi_t(x)$  can be thought of as a function of  $t$  with parameter  $x$  that also has a Gaussian-looking form.

## Another kernel

**Map :**  $\phi(x) = (1, \sqrt{2}x_1, \dots, \sqrt{2}x_d, x_1^2, \dots, x_d^2, \dots, \sqrt{2}x_i x_j, \dots)$

**Kernel :**  $\phi(x)^T \phi(x') = K(x, x') = (1 + x^T x')^2$

In fact, we can show  $K(x, x') = (1 + x^T x')^c$ , for  $c > 0$  is a kernel as well.

## Kernel arithmetic

Certain functions of kernels can produce new kernels.

Let  $K_1$  and  $K_2$  be any two kernels, then constructing  $K$  in the following ways produces a new kernel (among many other ways):

$$K(x, x') = K_1(x, x')K_2(x, x')$$

$$K(x, x') = K_1(x, x') + K_2(x, x')$$

$$K(x, x') = \exp\{K_1(x, x')\}$$

# KERNELIZED PERCEPTRON

## Returning to the Perceptron

We write the feature-expanded decision as

$$\begin{aligned}y_0 &= \text{sign} \left( \sum_{i \in \mathcal{S}} y_i \phi(x_0)^T \phi(x_i) \right) \\ &= \text{sign} \left( \sum_{i \in \mathcal{S}} y_i K(x_0, x_i) \right)\end{aligned}$$

We can pick the kernel we want to use. Let's pick the RBF (set  $a = 1$ ). Then

$$y_0 = \text{sign} \left( \sum_{i \in \mathcal{S}} y_i e^{-\frac{1}{b} \|x_0 - x_i\|^2} \right)$$

Notice that we never actually need to calculate  $\phi(x)$ .

What is this doing?

- ▶ Notice  $0 < K(x_0, x_i) \leq 1$ , with bigger values when  $x_0$  is closer to  $x_i$ .
- ▶ This is like a “soft voting” among the data picked by Perceptron.

# KERNELIZED PERCEPTRON

## Learning the kernelized Perceptron

Recall: Given a current vector  $w^{(t)} = \sum_{i \in \mathcal{S}_t} y_i x_i$ , we update it as follows,

1. Find a new  $x'$  such that  $y' \neq \text{sign}(x'^T w^{(t)})$
2. Add the index of  $x'$  to  $\mathcal{S}$  and set  $w^{(t+1)} = \sum_{i \in \mathcal{S}_{t+1}} y_i x_i$

Again we only need dot products, meaning these steps are equivalent to

1. Find a new  $x'$  such that  $y' \neq \text{sign}(\sum_{i \in \mathcal{S}_t} y_i K(x', x_i))$
2. Add the index of  $x'$  to  $\mathcal{S}$  *but don't bother calculating*  $w^{(t+1)}$

The trick is to realize that we never need to work with  $\phi(x)$ .

- ▶ We don't need  $\phi(x)$  to do Step 1 above.
- ▶ We don't need  $\phi(x)$  to classify new data (previous slide).
- ▶ We only ever need to calculate  $K(x, x')$  between two points.

# KERNEL $k$ -NN

## An extension

We can generalize kernelized Perceptron to *soft*  $k$ -NN with a simple change. Instead of summing over misclassified data  $S$ , sum over *all* the data:

$$y_0 = \text{sign} \left( \sum_{i=1}^n y_i e^{-\frac{1}{b} \|x_0 - x_i\|^2} \right).$$

Next, notice the *decision* doesn't change if we divide by a positive constant.

Let :  $Z = \sum_{j=1}^n e^{-\frac{1}{b} \|x_0 - x_j\|^2}$

Construct : Vector  $p(x_0)$ , where  $p_i(x_0) = \frac{1}{Z} e^{-\frac{1}{b} \|x_0 - x_i\|^2}$

Declare :  $y_0 = \text{sign} \left( \sum_{i=1}^n y_i p_i(x_0) \right)$

- ▶ We let all data vote for the label based on a “confidence score”  $p(x_0)$ .
- ▶ Set  $b$  so that most  $p_i(x_0) \approx 0$  to only focus on neighborhood around  $x_0$ .



# KERNEL REGRESSION

## Nadaraya-Watson model

The developments are almost limitless.

Here's a regression example almost identical to the kernelized  $k$ -NN:

Before:  $y \in \{-1, +1\}$

Now:  $y \in \mathbb{R}$

Using the RBF kernel, for a new  $(x_0, y_0)$  predict

$$y_0 = \sum_{i=1}^n y_i \frac{K(x_0, x_i)}{\sum_{j=1}^n K(x_0, x_j)}.$$

### What is this doing?

We're taking a locally weighted average of all  $y_i$  for which  $x_i$  is close to  $x_0$  (as decided by the kernel width). *Gaussian processes* are another option. . .

# GAUSSIAN PROCESSES (FOR REGRESSION)

# KERNELIZED BAYESIAN LINEAR REGRESSION

**Regression setup:** For  $n$  observations, with response vector  $y \in \mathbb{R}^n$  and their feature matrix  $X$ , we define the likelihood and prior

$$y \sim N(Xw, \sigma^2 I), \quad w \sim N(0, \lambda^{-1} I).$$

**Marginalizing:** What if we integrate out  $w$ ? We can solve this,

$$p(y|X) = \int p(y|X, w)p(w)dw = N(0, \sigma^2 I + \lambda^{-1} XX^T).$$

**Kernelization:** Notice that  $(XX^T)_{ij} = x_i^T x_j$ . Replace each  $x$  with  $\phi(x)$  after which we can say  $[\phi(X)\phi(X)^T]_{ij} = K(x_i, x_j)$ . We can define  $K$  directly, so

$$p(y|X) = \int p(y|X, w)p(w)dw = N(0, \sigma^2 I + \lambda^{-1} K).$$

This is called a *Gaussian process*. We never use  $w$  or  $\phi(x)$ , but just  $K(x_i, x_j)$ .

# GAUSSIAN PROCESSES

## Definition

- Let  $f(x) \in \mathbb{R}$  and  $x \in \mathcal{X}$ .
- Define the *kernel*  $K(x, x')$  between two points  $x$  and  $x'$ .
- Then  $f(x)$  is a *Gaussian process* and  $y(x)$  the noise-added process if for  $n$  observed pairs  $(x_1, y_1), \dots, (x_n, y_n)$ , where  $x \in \mathcal{X}$  and  $y \in \mathbb{R}$ ,


$$y | f \sim N(f, \sigma^2 I), \quad f \sim N(0, K) \quad \Longleftrightarrow \quad y \sim N(0, \sigma^2 I + K)$$

where  $y = (y_1, \dots, y_n)^T$  and  $K$  is  $n \times n$  with  $K_{ij} = K(x_i, x_j)$ .

## Comments:

- ▶ We assume  $\lambda = 1$  to reduce notation.
- ▶ Typical breakdown:  $f(x)$  is the GP and  $y(x)$  equals  $f(x)$  plus i.i.d. noise.
- ▶ The kernel is what keeps this from being “just a Gaussian.”


# GAUSSIAN PROCESSES



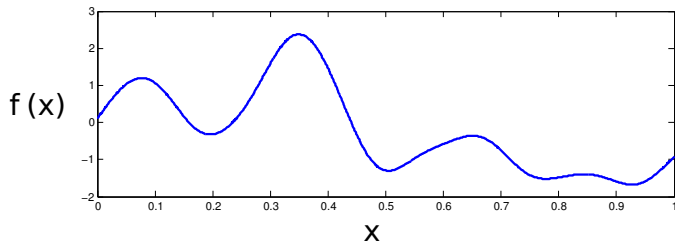
**Above:** A Gaussian process  $f(x)$  ( $x \in \mathbb{R}$ ) using

$$K(x_i, x_j) = \exp \left\{ -\frac{\|x_i - x_j\|^2}{b} \right\}.$$

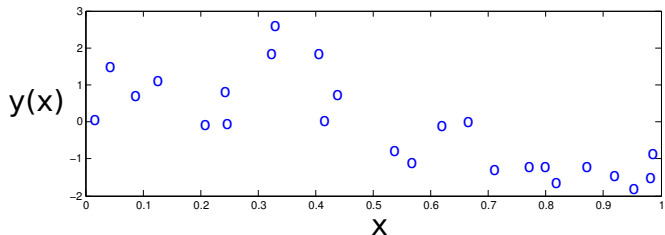
**Right:** The covariance of  $f(x)$  defined by  $K$ .



# GAUSSIAN PROCESSES



**Top:** Unobserved underlying function,  
**Bottom:** Noisy observed data sampled from this function



# PREDICTIONS WITH GAUSSIAN VECTORS

## Bayesian linear regression

Imagine we have  $n$  observation pairs  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$  and want to predict  $y_0$  given  $x_0$ . Integrating out  $w$  and setting  $\lambda = 1$ , the joint distribution is

$$\begin{bmatrix} y_0 \\ y \end{bmatrix} \sim \text{Normal} \left( \mathbf{0}, \sigma^2 I + \begin{bmatrix} x_0^T x_0 & (Xx_0)^T \\ Xx_0 & XX^T \end{bmatrix} \right)$$

We want to predict  $y_0$  given  $\mathcal{D}$  and  $x_0$ . Calculations can show that

$$\begin{aligned} y_0 | \mathcal{D}, x_0 &\sim \text{Normal}(\mu_0, \sigma_0^2) \\ \mu_0 &= (Xx_0)^T (\sigma^2 I + XX^T)^{-1} y \\ \sigma_0^2 &= \sigma^2 + x_0^T x_0 - (Xx_0)^T (\sigma^2 I + XX^T)^{-1} (Xx_0) \end{aligned}$$

Since the Gaussian process is only ever evaluated at a finite set of points, this is useful. Notice all the dot products between  $x_i$ . These can become kernels.

# PREDICTIONS WITH GAUSSIAN PROCESSES

## Predictive distribution of $y(x)$

Given measured data  $\mathcal{D}_n = \{(x_1, y_1), \dots, (x_n, y_n)\}$ , the distribution of  $y(x)$  can be calculated at any *new*  $x$  to make predictions.

Let  $K(x, \mathcal{D}_n) = [K(x, x_1), \dots, K(x, x_n)]$  and  $K_n$  be the  $n \times n$  kernel matrix restricted to points in  $\mathcal{D}_n$ . Then kernelizing the previous slide, we see that

$$y(x)|\mathcal{D}_n \sim N(\mu(x), \Sigma(x)),$$

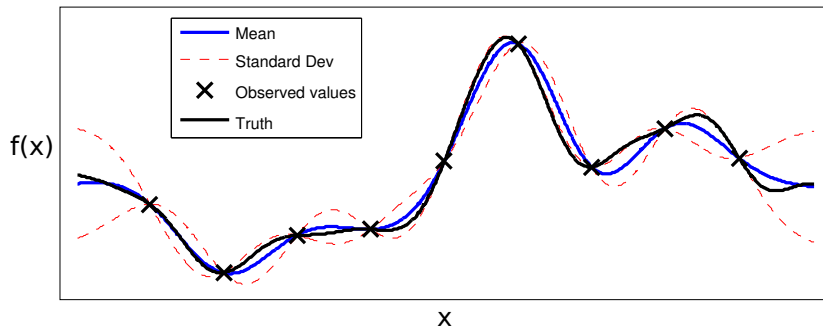
$$\mu(x) = K(x, \mathcal{D}_n)(\sigma^2 I + K_n)^{-1}y,$$

$$\Sigma(x) = \sigma^2 + K(x, x) - K(x, \mathcal{D}_n)(\sigma^2 I + K_n)^{-1}K(x, \mathcal{D}_n)^T$$

For the posterior of  $f(x)$  instead of  $y(x)$ , just remove  $\sigma^2$ .



# GAUSSIAN PROCESSES POSTERIOR



What does the posterior distribution of  $f(x)$  look like?

- ▶ We have data marked by an  $\times$ .
- ▶ These values pin down the function  $f(x)$  nearby
- ▶ We get a mean and variance for every possible  $x$  from previous slide.
- ▶ The distribution on  $y(x)$  adds variance  $\sigma^2$  (*very small above*) point-wise.