

COMS 4721: Machine Learning for Data Science

Lecture 13, 3/9/2021

Prof. John Paisley

Department of Electrical Engineering

Columbia University

BOOSTING

Robert E. Schapire and Yoav Freund, *Boosting: Foundations and Algorithms*, MIT Press, 2012.
See this textbook for many more details. (I borrow some figures from that book.)

BAGGING CLASSIFIERS

Algorithm: Bagging binary classifiers

Given $(x_1, y_1), \dots, (x_n, y_n), x \in \mathcal{X}, y \in \{-1, +1\}$

- ▶ For $b = 1, \dots, B$
 - ▶ Sample a bootstrap dataset \mathcal{B}_b of size n . For each entry in \mathcal{B}_b , select (x_i, y_i) with probability $\frac{1}{n}$. Some (x_i, y_i) will repeat and some won't appear in \mathcal{B}_b .
 - ▶ Learn a classifier f_b using data in \mathcal{B}_b .
- ▶ Define the classification rule to be

$$f_{\text{bag}}(x_0) = \text{sign} \left(\sum_{b=1}^B f_b(x_0) \right).$$

-
- ▶ With bagging, we observe that a *committee* of classifiers votes on a label.
 - ▶ Each classifier is learned on a *bootstrap sample* from the data set.
 - ▶ Learning a collection of classifiers is referred to as an *ensemble method*.

BOOSTING

How is it that a committee of blockheads can somehow arrive at highly reasoned decisions, despite the weak judgment of the individual members?

- Schapire & Freund, “Boosting: Foundations and Algorithms”

Boosting is another powerful method for ensemble learning. It is similar to bagging in that a set of classifiers are combined to make a better one.

It works for any classifier, but a “weak” one that is easy to learn is usually chosen. (weak = accuracy a little better than random guessing)

Short history

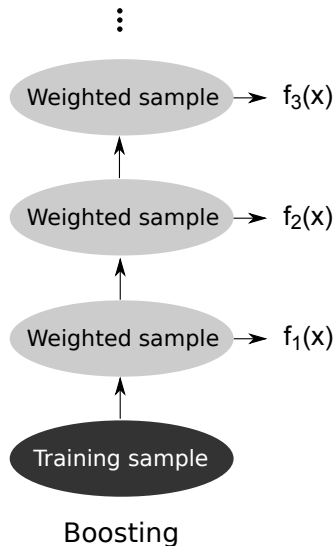
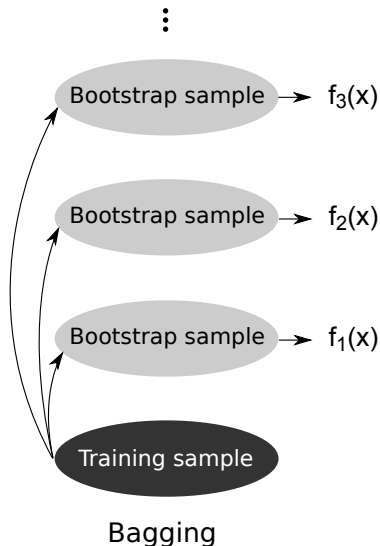
1984 : Leslie Valiant and Michael Kearns ask if “boosting” is possible.

1989 : Robert Schapire creates first boosting algorithm.

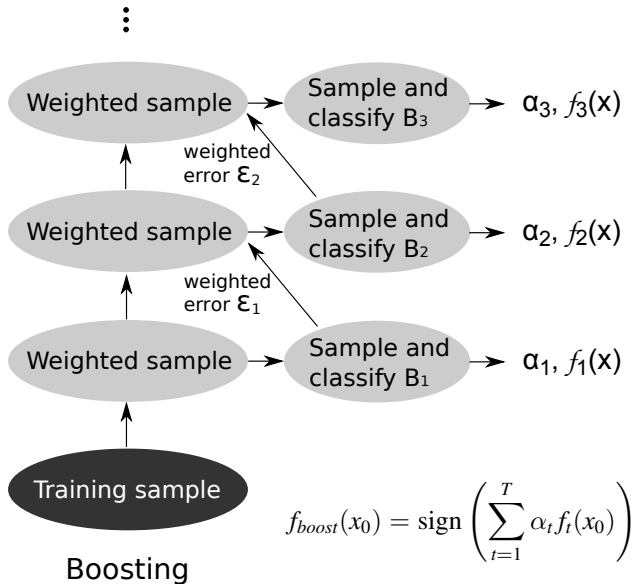
1990 : Yoav Freund creates an optimal boosting algorithm.

1995 : Freund and Schapire create AdaBoost (Adaptive Boosting), the major boosting algorithm.

BAGGING VS BOOSTING (OVERVIEW)



THE ADABOOST ALGORITHM (SAMPLING VERSION)



THE ADABOOST ALGORITHM (SAMPLING VERSION)

Algorithm: Boosting a binary classifier

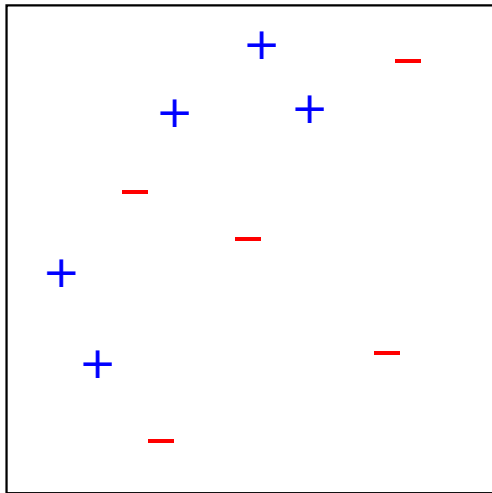
Given $(x_1, y_1), \dots, (x_n, y_n)$, $x \in \mathcal{X}$, $y \in \{-1, +1\}$, set $w_1(i) = \frac{1}{n}$ for $i = 1 : n$

- ▶ For $t = 1, \dots, T$
 1. Sample a new dataset \mathcal{B}_t of size n according to distribution w_t . Notice that we now pick (x_i, y_i) with probability $w_t(i)$ and not $\frac{1}{n}$.
 2. Learn a classifier f_t using data in \mathcal{B}_t .
 3. Set $\epsilon_t = \sum_{i=1}^n w_t(i) \mathbb{1}\{y_i \neq f_t(x_i)\}$ and $\alpha_t = \frac{1}{2} \ln \left(\frac{1-\epsilon_t}{\epsilon_t} \right)$.
 4. Scale $\hat{w}_{t+1}(i) = w_t(i)e^{-\alpha_t y_i f_t(x_i)}$ and set $w_{t+1}(i) = \frac{\hat{w}_{t+1}(i)}{\sum_j \hat{w}_{t+1}(j)}$.
- ▶ Set the classification rule to be

$$f_{\text{boost}}(x_0) = \text{sign} \left(\sum_{t=1}^T \alpha_t f_t(x_0) \right).$$

Comment: #1 & #2 usually simplified to “learn a classifier f_t using distribution w_t .”

BOOSTING A DECISION STUMP (EXAMPLE 1)

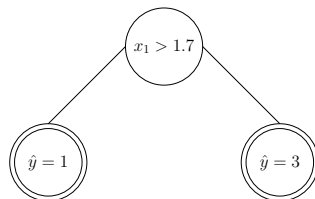


Original data


Uniform distribution, w_1

Learn *weak classifier*

Here: Use a decision stump



BOOSTING A DECISION STUMP (EXAMPLE 1)

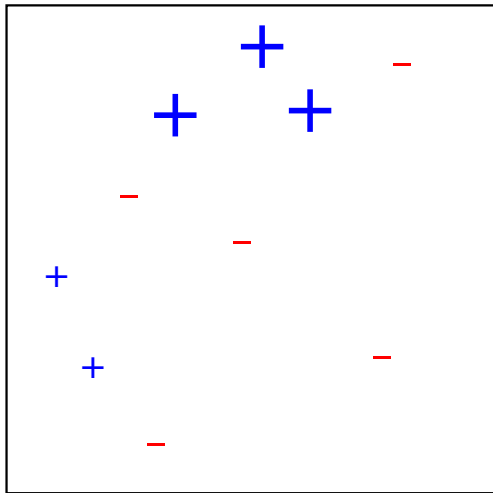


Round 1 classifier

Weighted error: $\epsilon_1 = 0.3$

Weight update: $\alpha_1 = 0.42$

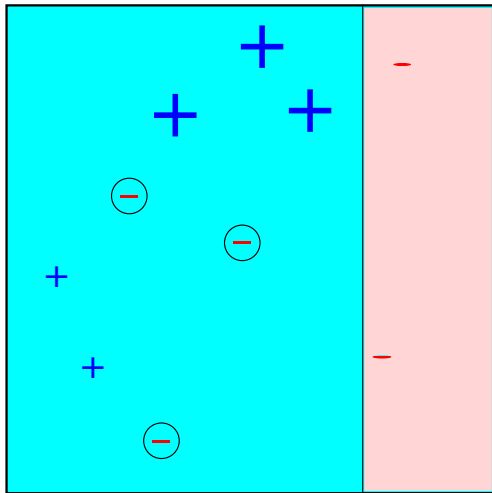
BOOSTING A DECISION STUMP (EXAMPLE 1)



Weighted data

After round 1

BOOSTING A DECISION STUMP (EXAMPLE 1)

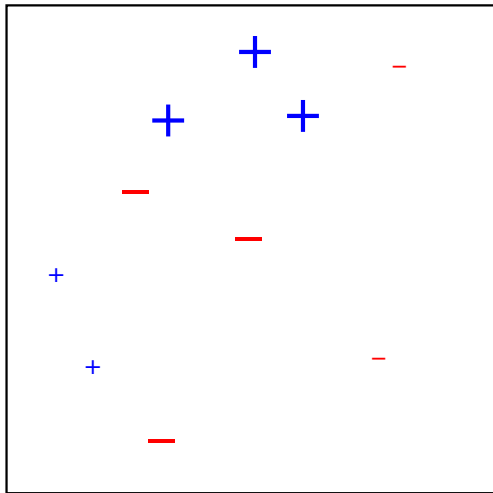


Round 2 classifier

Weighted error: $\epsilon_2 = 0.21$

Weight update: $\alpha_2 = 0.65$

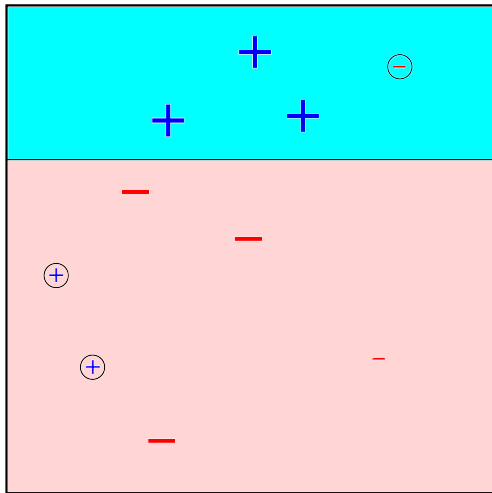
BOOSTING A DECISION STUMP (EXAMPLE 1)



Weighted data

After round 2

BOOSTING A DECISION STUMP (EXAMPLE 1)




Round 2 classifier

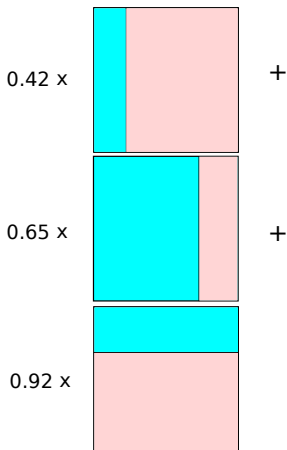
Weighted error: $\epsilon_3 = 0.14$

Weight update: $\alpha_3 = 0.92$


BOOSTING A DECISION STUMP (EXAMPLE 1)



Classifier after three rounds



BOOSTING A DECISION STUMP (EXAMPLE 2)



Example problem

Random guessing

50% error

Decision stump

45.8% error

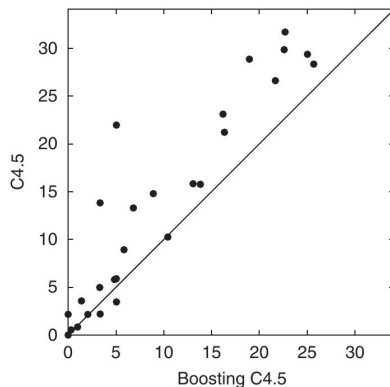
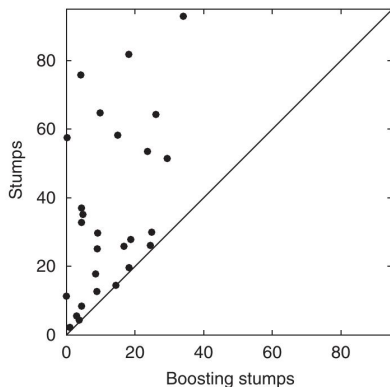
Full decision tree

24.7% error

Boosted stump

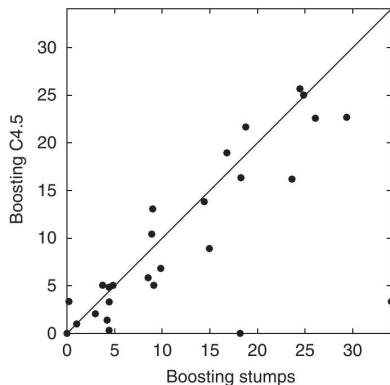
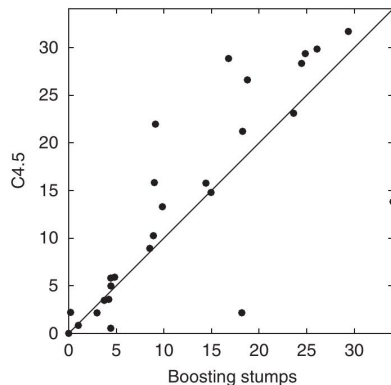
5.8% error

BOOSTING



Point = one dataset. Location = error rate w/ and w/o boosting. The boosted version of the same classifier almost always produces better results.

BOOSTING



(left) Boosting a bad classifier is often better than not boosting a good one.
(right) Boosting a good classifier is often better, but can take more time.

BOOSTING AND FEATURE MAPS

Q: What makes boosting work so well?

A: This is a well-studied question. We will present one analysis later, but we can also give intuition by tying it in with what we've already learned.

The classification for a new x_0 from boosting is

$$f_{boost}(x_0) = \text{sign} \left(\sum_{t=1}^T \alpha_t f_t(x_0) \right).$$

Define $\phi(x) = [f_1(x), \dots, f_T(x)]^\top$, where each $f_t(x) \in \{-1, +1\}$.

- ▶ We can think of $\phi(x)$ as a high dimensional feature map of x .
- ▶ The vector $\alpha = [\alpha_1, \dots, \alpha_T]^\top$ corresponds to a hyperplane.
- ▶ So the classifier can be written $f_{boost}(x_0) = \text{sign}(\phi(x_0)^\top \alpha)$.
- ▶ Boosting learns the feature mapping and hyperplane simultaneously.

APPLICATION: FACE DETECTION

FACE DETECTION (VIOLA & JONES, 2001)

Problem: Locate the faces in an image or video.

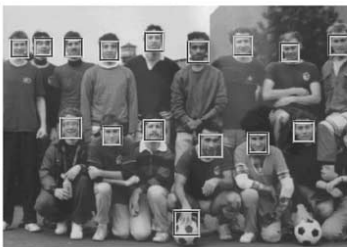
Processing: Divide image into patches of different scales, e.g., 24×24 , 48×48 , etc. Extract *features* from each patch.

Classify each patch as face or no face using a *boosted decision stump*. This can be done in real-time, for example by your digital camera (at 15 fps).



- ▶ One patch from a larger image. Mask it with many “feature extractors.”
- ▶ Each pattern gives one number, which is the sum of all pixels in black region minus sum of pixels in white region (total of 45,000+ features).

FACE DETECTION (EXAMPLE RESULTS)



ANALYSIS OF BOOSTING

ANALYSIS OF BOOSTING

Training error theorem

We can use *analysis* to make a statement about the accuracy of boosting *on the training data*.

Theorem: Under the AdaBoost framework, if ϵ_t is the weighted error of classifier f_t , then for the classifier $f_{boost}(x_0) = \text{sign}(\sum_{t=1}^T \alpha_t f_t(x_0))$,

$$\text{training error} = \frac{1}{n} \sum_{i=1}^n \mathbb{1}\{y_i \neq f_{boost}(x_i)\} \leq \exp\left(-2 \sum_{t=1}^T \left(\frac{1}{2} - \epsilon_t\right)^2\right).$$

Even if each ϵ_t is only a little better than random guessing, the sum over T classifiers can lead to a large negative value in the exponent when T is large.

For example, if we set:

$$\epsilon_t = 0.45, T = 1000 \rightarrow \text{training error} \leq 0.0067.$$

PROOF OF THEOREM

Setup

We break the proof into three steps. It is an application of the fact that

$$\text{if } \underbrace{a < b}_{\text{Step 2}} \quad \text{and} \quad \underbrace{b < c}_{\text{Step 3}} \quad \text{then} \quad \underbrace{a < c}_{\text{conclusion}}$$

- ▶ Step 1 calculates the value of b .
- ▶ Steps 2 and 3 prove the two inequalities.

Also recall the following step from AdaBoost:

- ▶ Update $\hat{w}_{t+1}(i) = w_t(i)e^{-\alpha_t y_i f_t(x_i)}$.
- ▶ Normalize $w_{t+1}(i) = \frac{\hat{w}_{t+1}(i)}{\sum_j \hat{w}_{t+1}(j)} \longrightarrow$ Define $Z_t = \sum_j \hat{w}_{t+1}(j)$.

PROOF OF THEOREM ($a \leq \mathbf{b} \leq c$)

Step 1

We first want to expand the equation of the weights to show that

$$w_{T+1}(i) = \frac{1}{n} \frac{e^{-y_i \sum_{t=1}^T \alpha_t f_t(x_i)}}{\prod_{t=1}^T Z_t} := \frac{1}{n} \frac{e^{-y_i h_T(x_i)}}{\prod_{t=1}^T Z_t} \rightarrow h_T(x_i) := \sum_{t=1}^T \alpha_t f_t(x_i)$$

Derivation of Step 1:

Notice the update rule: $w_{t+1}(i) = \frac{1}{Z_t} w_t(i) e^{-\alpha_t y_i f_t(x_i)}$

Do the same expansion for $w_t(i)$ and continue until reaching $w_1(i) = \frac{1}{n}$,

$$w_{T+1}(i) = w_1(i) \frac{e^{-\alpha_1 y_i f_1(x_i)}}{Z_1} \times \cdots \times \frac{e^{-\alpha_T y_i f_T(x_i)}}{Z_T}$$

The product $\prod_{t=1}^T Z_t$ is “ \mathbf{b} ” above. We use this form of $w_{T+1}(i)$ in Step 2.

PROOF OF THEOREM ($\mathbf{a} \leq \mathbf{b} \leq \mathbf{c}$)

Step 2

Next show the training error of $f_{boost}^{(T)}$ (boosting after T steps) is $\leq \prod_{t=1}^T Z_t$.
Currently we know

$$w_{T+1}(i) = \frac{1}{n} \frac{e^{-y_i h_T(x_i)}}{\prod_{t=1}^T Z_t} \Rightarrow w_{T+1}(i) \prod_{t=1}^T Z_t = \frac{1}{n} e^{-y_i h_T(x_i)} \quad \& \quad f_{boost}^{(T)}(x) = \text{sign}(h_T(x))$$

Derivation of Step 2:

Observe that $0 < e^{z_1}$ and $1 < e^{z_2}$ for any $z_1 < 0 < z_2$. Therefore

$$\underbrace{\frac{1}{n} \sum_{i=1}^n \mathbb{1}\{y_i \neq f_{boost}^{(T)}(x_i)\}}_a \leq \frac{1}{n} \sum_{i=1}^n e^{-y_i h_T(x_i)} \\ = \sum_{i=1}^n w_{T+1}(i) \prod_{t=1}^T Z_t = \underbrace{\prod_{t=1}^T Z_t}_b$$

“a” is the training error – the quantity we care about.

PROOF OF THEOREM $(a \leq \mathbf{b} \leq \mathbf{c})$

Step 3

The final step is to calculate an upper bound on Z_t , and by extension $\prod_{t=1}^T Z_t$.

Derivation of Step 3:

This step is slightly more involved. It also shows why $\alpha_t := \frac{1}{2} \ln \left(\frac{1-\epsilon_t}{\epsilon_t} \right)$.

$$\begin{aligned} Z_t &= \sum_{i=1}^n w_t(i) e^{-\alpha_t y_i f_t(x_i)} \\ &= \sum_{i: y_i = f_t(x_i)} e^{-\alpha_t} w_t(i) + \sum_{i: y_i \neq f_t(x_i)} e^{\alpha_t} w_t(i) \\ &= e^{-\alpha_t} (1 - \epsilon_t) + e^{\alpha_t} \epsilon_t \end{aligned}$$

Remember we defined $\epsilon_t = \sum_{i: y_i \neq f_t(x_i)} w_t(i)$, the probability of error for w_t .

PROOF OF THEOREM ($a \leq \mathbf{b} \leq \mathbf{c}$)

Derivation of Step 3 (continued):

Remember from Step 2 that

$$\text{training error} = \frac{1}{n} \sum_{i=1}^n \mathbb{1}\{y_i \neq f_{\text{boost}}(x_i)\} \leq \prod_{t=1}^T Z_t.$$

and we just showed that $Z_t = e^{-\alpha_t}(1 - \epsilon_t) + e^{\alpha_t}\epsilon_t$.

We want the training error to be small, so we pick α_t to *minimize* Z_t . Minimizing, we get the value of α_t used by AdaBoost:

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right).$$

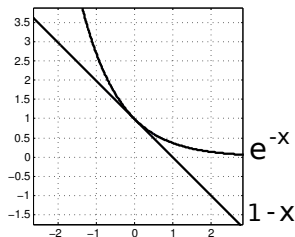
Plugging this value back in gives $Z_t = 2\sqrt{\epsilon_t(1 - \epsilon_t)}$.

PROOF OF THEOREM ($a \leq b \leq c$)

Derivation of Step 3 (continued):

Next, re-write Z_t as

$$\begin{aligned} Z_t &= 2\sqrt{\epsilon_t(1-\epsilon_t)} \\ &= \sqrt{1-4\left(\frac{1}{2}-\epsilon_t\right)^2} \end{aligned}$$



Then, use the inequality $1 - x \leq e^{-x}$ to conclude that

$$Z_t = \left(1 - 4\left(\frac{1}{2} - \epsilon_t\right)^2\right)^{\frac{1}{2}} \leq \left(e^{-4\left(\frac{1}{2} - \epsilon_t\right)^2}\right)^{\frac{1}{2}} = e^{-2\left(\frac{1}{2} - \epsilon_t\right)^2}.$$

PROOF OF THEOREM

Concluding the right inequality ($a \leq \mathbf{b} \leq c$)

Because both sides of $Z_t \leq e^{-2(\frac{1}{2}-\epsilon_t)^2}$ are positive, we can say that

$$\prod_{t=1}^T Z_t \leq \prod_{t=1}^T e^{-2(\frac{1}{2}-\epsilon_t)^2} = e^{-2 \sum_{t=1}^T (\frac{1}{2}-\epsilon_t)^2}.$$

This concludes the “ $b \leq c$ ” portion of the proof.

Combining everything

$$\text{training error} = \overbrace{\frac{1}{n} \sum_{i=1}^n \mathbb{1}\{y_i \neq f_{\text{boost}}(x_i)\}}^a \leq \overbrace{\prod_{t=1}^T Z_t}^b \leq \overbrace{e^{-2 \sum_{t=1}^T (\frac{1}{2}-\epsilon_t)^2}}^c.$$

We set out to prove “ $a \leq c$ ” and we did so by using “ b ” as a stepping-stone.

TRAINING VS TESTING ERROR

Q: Driving the training error to zero leads one to ask, does boosting overfit?

A: Sometimes, but very often it doesn't!

